


# LoRA: Low-Rank Adaptation of Large Language Models

## Motivation: The Fine-tuning Problem



Fine-tuning large language models presents significant challenges:

- **Parameter Inefficiency:** Fine-tuning requires updating all parameters (often billions)
- **Memory Requirements:** Storing optimizer states, gradients, and model copies demands substantial VRAM
- **Storage Overhead:** Each fine-tuned model requires saving a complete copy of all parameters
- **Training Cost:** Full fine-tuning is computationally expensive and time-consuming

## Enter LoRA: Parameter-Efficient Fine-tuning

LoRA (Low-Rank Adaptation) solves these challenges by recognizing that updates to pre-trained weights during fine-tuning have a **low intrinsic rank**. By low intrinsic rank we mean that when a large pre-trained model is fine-tuned for a specific task, the changes required in its weight matrices are structured and do not span the full high-dimensional space of the original weights. Instead of updating all weights directly, LoRA injects trainable rank decomposition matrices into the model.

# Core Concept: Low-Rank Decomposition

The key insight is that we can represent the weight update matrix as a product of two lower-rank matrices:

$$\Delta W = B \times A$$

Where:

- Original weight matrix  $W$  has dimensions  $d \times k$
- $B$  has dimensions  $d \times r$
- $A$  has dimensions  $r \times k$
- $r \ll \min(d, k)$  (significantly lower rank than the original matrix)

The updated forward pass becomes:

$$h = Wx + \Delta Wx = Wx + BAx$$

## LoRA in Practice: Applying to Transformer Weights

LoRA typically targets specific weight matrices in the model:

- **Query and Key projections** in self-attention layers
- **Value projections** in self-attention layers
- **Output projections** after self-attention
- **Feed-forward network** weights

For each target weight matrix, we:

1. Freeze the original pre-trained weights  $W$
2. Initialize low-rank matrices  $A$  and  $B$ 
  - $A$  is initialized with random Gaussian values
  - $B$  is initialized with zeros (ensuring no impact at the start)
3. Scale the contribution with a parameter  $\alpha$
4. Apply the update during the forward pass:  $Wx + \frac{\alpha}{r}(BA)x$

## Advantages of LoRA

- **Parameter Efficiency:** Only trains a small fraction of parameters (typically <1%)
- **Memory Efficiency:** Significantly reduced memory footprint during training
- **Storage Efficiency:** Can store multiple task adaptations with minimal overhead

- **Inference Efficiency:** LoRA matrices can be merged with original weights at inference time
- **Composition:** Multiple LoRA adaptations can be combined or swapped

# Mathematical Representation

## Standard Fine-tuning

During traditional fine-tuning, we update weights from  $W$  to  $W'$ :

$$W' = W + \Delta W$$

## LoRA Approach

With LoRA, we approximate the update using low-rank matrices:

$$W' = W + \frac{\alpha}{r}BA$$

Where:

- $W \in \mathbb{R}^{d \times k}$  is the pre-trained weight matrix
- $B \in \mathbb{R}^{d \times r}$  and  $A \in \mathbb{R}^{r \times k}$  are trainable low-rank matrices
- $r$  is the rank (hyperparameter, typically 4-64)
- $\alpha$  is a scaling factor (hyperparameter)

## Training Process

1. Initialize  $A$  with random Gaussian values:  $A \sim N(0, \sigma^2)$
2. Initialize  $B$  with zeros:  $B = 0$
3. Freeze the pre-trained weights  $W$
4. Train only the parameters in  $A$  and  $B$
5. During inference, compute:  $h = Wx + \frac{\alpha}{r}(BA)x$

## Hyperparameters

The most critical hyperparameters in LoRA are:

- **Rank ( $r$ ):** Controls the expressiveness of the adaptation (higher values = more capacity)
- **Alpha ( $\alpha$ ):** Scales the contribution of the low-rank update
- **Target modules:** Which weight matrices to apply LoRA to

- **Learning rate:** Often higher learning rates can be used compared to full fine-tuning

# Performance Comparison

Method	Trainable Parameters	Memory Usage	Storage	Training Time
Full Fine-tuning	100%	High	Full model copy	Slow
LoRA (r=8)	<1%	Low	Small adapter	Fast
LoRA (r=16)	<1%	Low	Small adapter	Fast
LoRA (r=32)	<1%	Low	Small adapter	Fast

# In This Assignment

You will implement LoRA for your SmolLM model, focusing on:

1. Creating the LoRA adaptation layers
2. Freezing pre-trained weights
3. Applying LoRA to specific weight matrices
4. Training and evaluating the LoRA-adapted model
5. Comparing performance against full fine-tuning

# Additional Resources

## Original Paper

- [LoRA: Low-Rank Adaptation of Large Language Models](#) - Hu et al., 2021

## Video Explanations

- [LoRA explained \(and a bit about precision and quantization\)](#) - Good introduction to LoRA
- [Fine-tuning LLMs with PEFT and LoRA](#) - Practical walkthrough